

dslusser_3

David Slusser

10/18/2020

R Markdown

This is homework assignment number 3, using naive bayes to predict if a flight will be on time or not

```
library(readr) # Need to load data
library(dplyr) # Selecting variables
```

```
## Warning: package 'dplyr' was built under R version 4.0.2
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(caret) # For splitting into training and validation
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
## Warning: package 'ggplot2' was built under R version 4.0.2
```

```
library(ISLR)
```

```
## Warning: package 'ISLR' was built under R version 4.0.2
```

```
library(e1071) # For the Naive Bayes model
```

```
## Warning: package 'e1071' was built under R version 4.0.2
```

```
library(gmodels) # For counts table
```

```
## Warning: package 'gmodels' was built under R version 4.0.2
```

```
library(pROC) # For the ROC plot
```

```
## Type 'citation("pROC")' for a citation.
```

```
##  
## Attaching package: 'pROC'
```

```
## The following object is masked from 'package:gmodels':  
##  
##      ci
```

```
## The following objects are masked from 'package:stats':  
##  
##      cov, smooth, var
```

```
flights <- read_csv("~/Desktop/School/Graduate/Machine Learning/HW/Data/FlightDelays.csv")
```

```
## Parsed with column specification:  
## cols(  
##   CRS_DEP_TIME = col_double(),  
##   CARRIER = col_character(),  
##   DEP_TIME = col_double(),  
##   DEST = col_character(),  
##   DISTANCE = col_double(),  
##   FL_DATE = col_character(),  
##   FL_NUM = col_double(),  
##   ORIGIN = col_character(),  
##   Weather = col_double(),  
##   DAY_WEEK = col_double(),  
##   DAY_OF_MONTH = col_double(),  
##   TAIL_NUM = col_character(),  
##   `Flight Status` = col_character()  
## )
```

```
# We need to make variables factors
flights$DAY_WEEK <- cut(flights$DAY_WEEK,c(-Inf, 1, 2, 3, 4, 5, 6, Inf), # Convert day number to days
                        labels=c("Monday", "Tuesday", "Wednesday", "Thursday", "Friday", "Saturday", "Sunday"))

flightsModel <- flights %>%
  select(CRS_DEP_TIME, CARRIER, DEST, ORIGIN, DAY_WEEK, `Flight Status`) %>%
  rename(Status = `Flight Status`)

# Convert departure time into numeric
flightsModel$CRS_DEP_TIME <- as.factor(flightsModel$CRS_DEP_TIME)

# We want the count and proportions for each airport for delayed
flightsModel %>%
  select(Status, ORIGIN) %>% # Select the variables needed
  group_by(ORIGIN, Status) %>% # Group by origin and then status to get the amount delayed/on time for airport
  summarise(Count = n()) %>% # Count the number of flights delayed/ontime at each airport
  mutate(Freq = Count / sum(Count)) # Get the frequency of flights delayed/ontime at each airport
```

```
## `summarise()` regrouping output by 'ORIGIN' (override with `.groups` argument)
```

ORIGIN <chr>	Status <chr>	Count <int>	Freq <dbl>
BWI	delayed	37	0.2551724
BWI	ontime	108	0.7448276
DCA	delayed	221	0.1613139
DCA	ontime	1149	0.8386861
IAD	delayed	170	0.2478134
IAD	ontime	516	0.7521866
6 rows			

At BWI (Baltimore-Washington), there were 145 flights with 37 (25.5%) delayed and 108 (74.5%) on time At DCA (Reagan National), there were 1370 flights with 221 (16.1%) delayed and 1149 (83.9%) on time At IAD (Dulles), there were 686 flights with 170 (24.8%) delayed and 516 (75.2%) on time

Overall, there were 2,201 flights with 428 (19.4%) delayed and 1773 (80.6%) on time

```
# We need to create and build the naive bayes model
# Start with training and validation

# Set seed and divide the data into training (60%) and validation (40%)
set.seed(1234)
Index_Train <- createDataPartition(flightsModel$Status, p= 0.6, list = FALSE)
Train <- flightsModel[Index_Train,]
```

```
## Warning: The `i` argument of ``[`()`` can't be a matrix as of tibble 3.0.0.
## Convert to a vector.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_warnings()` to see where this warning was generated.
```

```
Validation <- flightsModel[-Index_Train,]

# Build a naive bayes classifier
nb_model <-naiveBayes(as.factor(Status) ~ as.factor(CRS_DEP_TIME) + CARRIER + DEST + ORI
GIN + DAY_WEEK,
                    data = Train) # This is the classifier with our categorical variab
les
nb_model # Produces the A-prori probabilities (Probabilities from deductive reasoning)
```

```

##
## Naive Bayes Classifier for Discrete Predictors
##
## Call:
## naiveBayes.default(x = X, y = Y, laplace = laplace)
##
## A-priori probabilities:
## Y
##   delayed   ontime
## 0.1945496 0.8054504
##
## Conditional probabilities:
##       as.factor(CRS_DEP_TIME)
## Y           600           630           640           645           700
## delayed 0.0038910506 0.0038910506 0.0194552529 0.0038910506 0.0350194553
## ontime  0.0140977444 0.0300751880 0.0075187970 0.0140977444 0.0404135338
##       as.factor(CRS_DEP_TIME)
## Y           730           735           759           800           830
## delayed 0.0077821012 0.0077821012 0.0000000000 0.0233463035 0.0077821012
## ontime  0.0122180451 0.0065789474 0.0009398496 0.0159774436 0.0150375940
##       as.factor(CRS_DEP_TIME)
## Y           840           845           850           900           925
## delayed 0.0272373541 0.0000000000 0.0116731518 0.0311284047 0.0000000000
## ontime  0.0244360902 0.0009398496 0.0150375940 0.0357142857 0.0018796992
##       as.factor(CRS_DEP_TIME)
## Y           930           1000          1030          1039          1040
## delayed 0.0038910506 0.0000000000 0.0272373541 0.0000000000 0.0038910506
## ontime  0.0225563910 0.0122180451 0.0300751880 0.0028195489 0.0065789474
##       as.factor(CRS_DEP_TIME)
## Y           1100          1130          1200          1230          1240
## delayed 0.0116731518 0.0038910506 0.0000000000 0.0038910506 0.0194552529
## ontime  0.0244360902 0.0122180451 0.0150375940 0.0159774436 0.0159774436
##       as.factor(CRS_DEP_TIME)
## Y           1245          1300          1315          1330          1359
## delayed 0.0466926070 0.0272373541 0.0038910506 0.0000000000 0.0116731518
## ontime  0.0291353383 0.0545112782 0.0018796992 0.0084586466 0.0131578947
##       as.factor(CRS_DEP_TIME)
## Y           1400          1430          1455          1500          1515
## delayed 0.0116731518 0.0194552529 0.1011673152 0.0350194553 0.0077821012
## ontime  0.0225563910 0.0206766917 0.0545112782 0.0328947368 0.0018796992
##       as.factor(CRS_DEP_TIME)
## Y           1520          1525          1530          1600          1605
## delayed 0.0000000000 0.0233463035 0.0233463035 0.0233463035 0.0000000000
## ontime  0.0000000000 0.0065789474 0.0216165414 0.0169172932 0.0000000000
##       as.factor(CRS_DEP_TIME)
## Y           1610          1630          1640          1645          1700
## delayed 0.0038910506 0.0233463035 0.0116731518 0.0038910506 0.0233463035
## ontime  0.0150375940 0.0281954887 0.0112781955 0.0159774436 0.0310150376
##       as.factor(CRS_DEP_TIME)
## Y           1710          1715          1720          1725          1730
## delayed 0.0116731518 0.0505836576 0.0311284047 0.0000000000 0.0194552529
## ontime  0.0112781955 0.0197368421 0.0075187970 0.0009398496 0.0159774436
##       as.factor(CRS_DEP_TIME)

```

```
## Y          1800          1830          1900          1930          2000
## delayed 0.0038910506 0.0350194553 0.0700389105 0.0116731518 0.0155642023
## ontime  0.0150375940 0.0272556391 0.0366541353 0.0093984962 0.0140977444
##          as.factor(CRS_DEP_TIME)
## Y          2030          2100          2120          2130
## delayed 0.0116731518 0.0116731518 0.0700389105 0.0000000000
## ontime  0.0112781955 0.0178571429 0.0328947368 0.0009398496
##
##          CARRIER
## Y          CO          DH          DL          MQ          OH          RU
## delayed 0.03501946 0.33852140 0.12451362 0.15953307 0.01167315 0.23735409
## ontime  0.03571429 0.22744361 0.19924812 0.11842105 0.01691729 0.18984962
##          CARRIER
## Y          UA          US
## delayed 0.01167315 0.08171206
## ontime  0.01503759 0.19736842
##
##          DEST
## Y          EWR          JFK          LGA
## delayed 0.3657588 0.2178988 0.4163424
## ontime  0.2913534 0.1682331 0.5404135
##
##          ORIGIN
## Y          BWI          DCA          IAD
## delayed 0.08560311 0.48249027 0.43190661
## ontime  0.06484962 0.64473684 0.29041353
##
##          DAY_WEEK
## Y          Monday    Tuesday    Wednesday    Thursday    Friday    Saturday
## delayed 0.20622568 0.13229572 0.13618677 0.14396887 0.17509728 0.04280156
## ontime  0.13721805 0.12593985 0.14379699 0.15601504 0.19360902 0.13157895
##          DAY_WEEK
## Y          Sunday
## delayed 0.16342412
## ontime  0.11184211
```

```
# Predict the status of the flight using the model and the validation data set
predicted_status <- predict(nb_model, Validation) # this provides the predicted label (de
layed/on time) where
```

```
# P > 0.5 is the cutoff
```

```
# Show the confusion matrix of the flight status
```

```
CrossTable(x = Validation$Status, y = predicted_status, prop.chisq = FALSE)
```

```
##
##
##      Cell Contents
## |-----|
## |                      N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  880
##
##
##      predicted_status
## Validation$Status |   delayed |   ontime | Row Total |
## -----|-----|-----|-----|
##           delayed |         7 |        164 |        171 |
##                   |    0.041 |    0.959 |    0.194 |
##                   |    0.259 |    0.192 |           |
##                   |    0.008 |    0.186 |           |
## -----|-----|-----|-----|
##           ontime  |        20 |        689 |        709 |
##                   |    0.028 |    0.972 |    0.806 |
##                   |    0.741 |    0.808 |           |
##                   |    0.023 |    0.783 |           |
## -----|-----|-----|-----|
##      Column Total |        27 |        853 |        880 |
##                   |    0.031 |    0.969 |           |
## -----|-----|-----|-----|
##
##
```

For the PROC we want the probabilities of each, so the predicted status we need the raw probabilities

```
predicted_status_prob <-predict(nb_model, Validation, type = "raw") # type = "raw" provides prob of each
```

```
head(predicted_status_prob) # gives first couple rows of the predicted prob. delayed is column 1 and
```

```
##      delayed   ontime
## [1,] 0.2082043 0.7917957
## [2,] 0.2433158 0.7566842
## [3,] 0.2754174 0.7245826
## [4,] 0.2754174 0.7245826
## [5,] 0.3898837 0.6101163
## [6,] 0.3898837 0.6101163
```

```
# on time is column 2
```

```
# Get the ROC curve  
# This uses the pROC package  
roc(Validation$Status, predicted_status_prob[,2]) # Prob that the flight is ontime
```


```
## Setting levels: control = delayed, case = ontime
```

```
## Setting direction: controls < cases
```

```
##  
## Call:  
## roc.default(response = Validation$Status, predictor = predicted_status_prob[, 2])  
##  
## Data: predicted_status_prob[, 2] in 171 controls (Validation$Status delayed) < 709 cases (Validation$Status ontime).  
## Area under the curve: 0.626
```

```
plot.roc(Validation$Status, predicted_status_prob[,2]) # Plot of curve that the flight is ontime
```

```
## Setting levels: control = delayed, case = ontime  
## Setting direction: controls < cases
```

This cross table shows us that we have 184 misclassifications

We find the Area under the curve (AUC) is 62.6%

Better curves are closer to the top left corner, but as our ROC is closer to the dashed line, our model is not as strong (also seen with the lower AUC)